

# Exploring the Existing Namespaces

In this lesson, we will explore and discuss briefly about the existing Namespaces.

## WE'LL COVER THE FOLLOWING ^

- Getting the Existing Namespaces
- The default Namespace
- The kube-public Namespace
- The kube-system Namespace

Now that we know that our cluster has multiple Namespaces, let's explore them a bit.

## Getting the Existing Namespaces #

We can list all the Namespaces through the `kubectl get namespaces` command. As with the most of the other Kubernetes objects and resources, we can also use a shortcut `ns` instead of the full name.

```
kubectl get ns
```



The **output** is as follows.

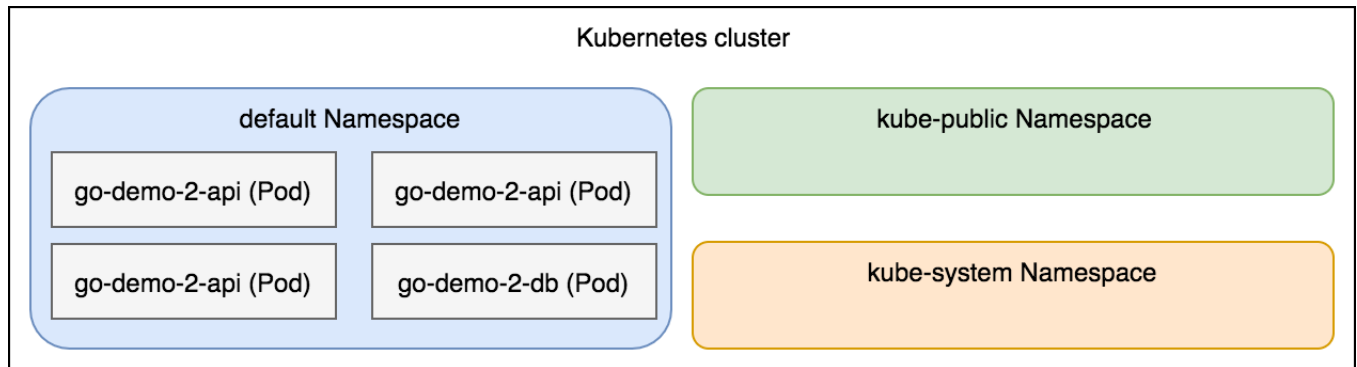
NAME	STATUS	AGE
default	Active	23m
kube-node-lease	Active	23m
kube-public	Active	23m
kube-system	Active	23m



We can see that three Namespaces were set up automatically when we created the Minikube cluster.

## The default Namespace #

The `default` Namespace is the one we used all this time. If we do not specify otherwise, all the `kubectl` commands will operate against the objects in the `default` Namespace. That's where our `go-demo-2` application is running. Even though we were not aware of its existence, we now know that's where the objects we created are placed.



The Namespaces and the go-demo-2 Pods

There are quite a few ways to specify a Namespace. For now, we'll use the `--namespace` argument. It is one of the global options that is available for all `kubectl` commands.

## The kube-public Namespace #

The command that will retrieve all the objects from the `kube-public` Namespace is as follows.

```
kubectl --namespace kube-public get all
```



The **output** states that `No resources were found`. That's disappointing, isn't it? Kubernetes does not use the `kube-public` Namespace for its system-level objects. All the objects we created are in the `default` Namespace.

The `kube-public` Namespace is readable by all users from all Namespaces.

The primary reason for `kube-public`'s existence is to provide space where we can create objects that should be visible throughout the whole cluster.

A good example is ConfigMaps. When we create one in, let's say, the `default` Namespace, it is accessible only by the other objects in the same Namespace. Those residing somewhere else would be oblivious of its existence. If we'd like such a ConfigMap to be visible to all objects no matter where they are, we'd put it into the `kube-public` Namespace instead. We won't use this Namespace much (if at all).

## The kube-system Namespace #

The `kube-system` Namespace is critical.

Almost all the objects and resources Kubernetes needs are running inside `kube-system` Namespace.

We can check that by executing the command that follows.

```
kubectl --namespace kube-system get all
```



We retrieved all the objects and resources running inside the `kube-system` Namespace. The **output** is as follows.

NAME	READY	STATUS	RESTARTS	AGE
pod/coredns-fb8b8dccf-54ppw	1/1	Running	1	4h32m
pod/coredns-fb8b8dccf-vgxjk	1/1	Running	1	4h32m
pod/default-http-backend-6864bbb7db-7mtdc	1/1	Running	0	4h32m
pod/etcd-minikube	1/1	Running	0	4h31m
pod/kube-addon-manager-minikube	1/1	Running	0	4h31m
pod/kube-apiserver-minikube	1/1	Running	0	4h31m
pod/kube-controller-manager-minikube	1/1	Running	0	4h31m
pod/kube-proxy-6w4xt	1/1	Running	0	4h32m
pod/kube-scheduler-minikube	1/1	Running	0	4h31m
pod/kubernetes-dashboard-79dd6bfc48-ztpd6	1/1	Running	4	4h32m
pod/nginx-ingress-controller-586cdc477c-r9b2m	1/1	Running	0	4h32m
pod/storage-provisioner	1/1	Running	0	4h32m

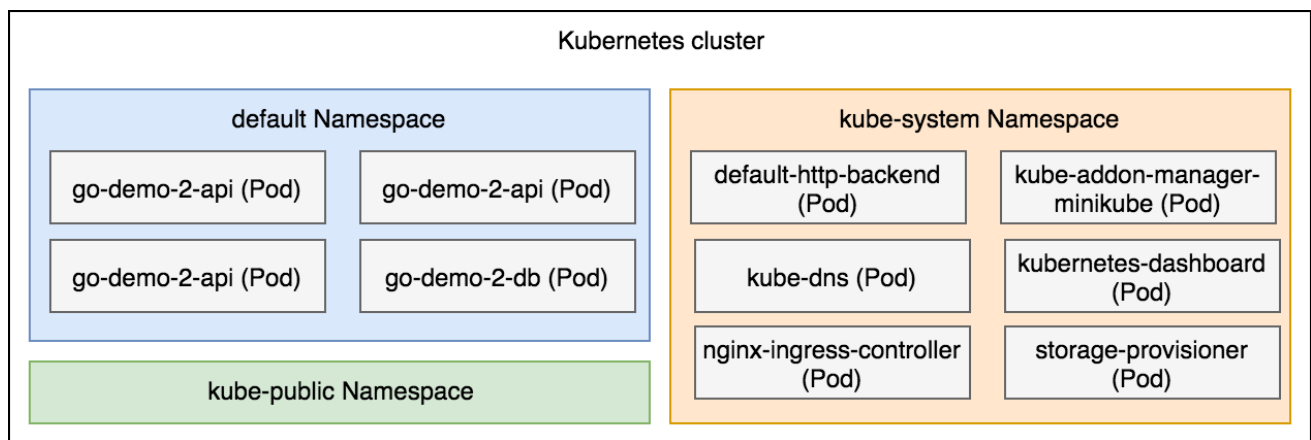
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
service/default-http-backend	NodePort	10.102.50.182	<none>	80:30001/TCP
service/kube-dns	ClusterIP	10.96.0.10	<none>	53/UDP,53/TCP,9153/TCP
service/kubernetes-dashboard	ClusterIP	10.102.236.69	<none>	80/TCP

NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR
daemonset.apps/kube-proxy	1	1	1	1	1	<none>

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/coredns	2/2	2	2	4h32m
deployment.apps/default-http-backend	1/1	1	1	4h32m
deployment.apps/kubernetes-dashboard	1/1	1	1	4h32m

deployment.apps/nginx-ingress-controller	1/1	1	1	4h32m	
NAME		DESIRED	CURRENT	READY	AGE
replicaset.apps/coredns-fb8b8dccf		2	2	2	4h32m
replicaset.apps/default-http-backend-6864bbb7db		1	1	1	4h32m
replicaset.apps/kubernetes-dashboard-79dd6bfc48		1	1	1	4h32m
replicaset.apps/nginx-ingress-controller-586cdc477c		1	1	1	4h32m

As we can see, quite a few things are running inside the `kube-system` Namespace. For example, we knew that there is an nginx Ingress controller, but this is the first time we saw its objects. It consists of a Replication Controller `nginx-ingress-controller`, and the Pod it created, `nginx-ingress-controller-fxrhn`.



The Namespaces and the Pods

As long as the system works as expected, there isn't much need to do anything inside the `kube-system` Namespace. The real fun starts when we create new Namespaces.

In the next lesson, we will create a Namespace and deploy to it.