

Sequential Breakdown and Verification of the Update Process

In this lesson, we will go through the sequential breakdown of the cluster update process and verify the update to the cluster.

WE'LL COVER THE FOLLOWING ^

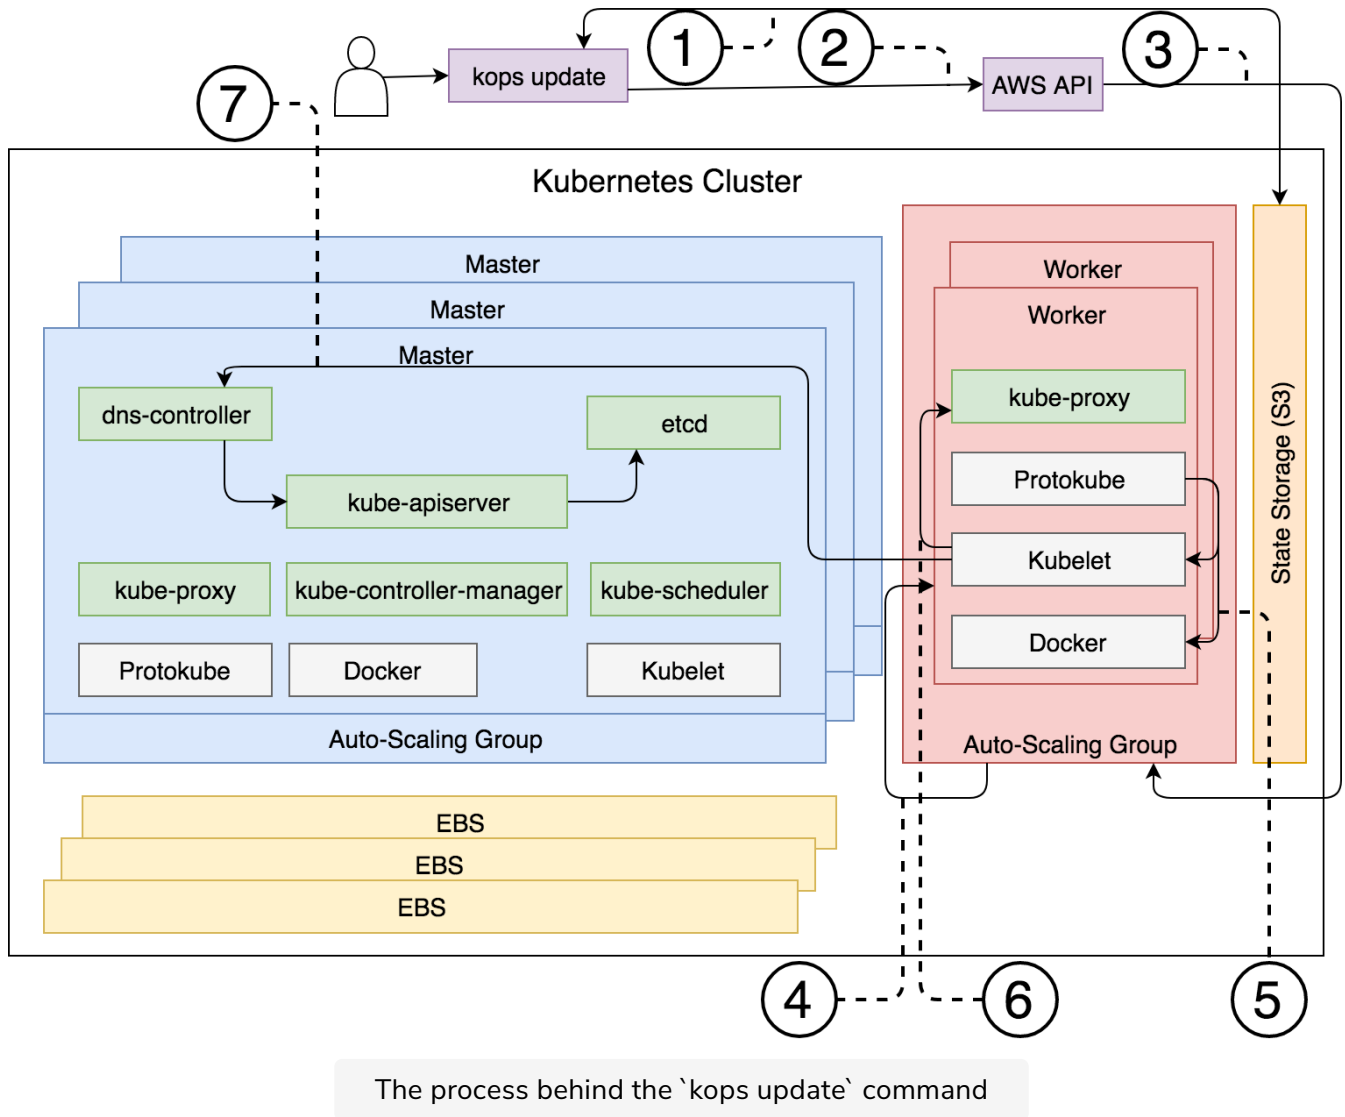
- Sequential Breakdown of the Update Process
- Verification of the Update Process

Sequential Breakdown of the Update Process

Let's see what happened when we executed the `kops update` command.

1. Kops retrieved the desired state from the S3 bucket.
2. Kops sent requests to AWS API to change the values of the workers ASG.
3. AWS modified the values of the workers ASG by increasing them by 1.
4. ASG created a new EC2 instance to comply with the new sizing.
5. Protokube installed Kubelet and Docker and created the manifest file with the list of Pods.
6. Kubelet read the manifest file and run the container that forms the `kube-proxy` Pod (the only Pod on the worker nodes).
7. Kubelet sent a request to the `kube-apiserver` (through the `dns-controller`) to register the new node and join it to the cluster. The information about the new node is stored in `etcd`.

This process is almost identical to the one used to create the nodes of the cluster.



The above illustration shows the sequence of the steps involved in the cluster update process.

Verification of the Update Process

Unless you are a very fast reader, ASG already created a new EC2 instance, and Kubelet joined it to the cluster. We can confirm that through the `kops validate` command.

```
kops validate cluster
```

The **output** is as follows.

```
Validating cluster devops23.k8s.local
```

```
INSTANCE GROUPS
```

```
NAME          ROLE    MACHINETYPE  MIN  MAX  SUBNETS
```

```
master-us-east-2a Master  t2.small     1    1    us-east-2a
```

```
master-us-east-2a Master t2.small 1 1 us-east-2a
master-us-east-2b Master t2.small 1 1 us-east-2b
master-us-east-2c Master t2.small 1 1 us-east-2c
nodes Node t2.small 2 2 us-east-2a,us-east-2b,us-east-2c
```

NODE STATUS

NAME	ROLE	READY
ip-172-20-120-133...	master	True
ip-172-20-33-237...	node	True
ip-172-20-34-249...	master	True
ip-172-20-65-28...	master	True
ip-172-20-95-101...	node	True

Your cluster devops23.k8s.local is ready

We can see that now we have two nodes (there was one before) and that they are located somewhere inside the three `us-east-2` availability zones.

Similarly, we can use `kubectl` to confirm that Kubernetes indeed added the new worker node to the cluster.

```
kubectl get nodes
```



The **output** is as follows.

NAME	STATUS	ROLES	AGE	VERSION
ip-172-20-120-133...	Ready	master	13m	v1.9.1
ip-172-20-33-237...	Ready	node	1m	v1.9.1
ip-172-20-34-249...	Ready	master	13m	v1.9.1
ip-172-20-65-28...	Ready	master	13m	v1.9.1
ip-172-20-95-101...	Ready	node	12m	v1.9.1



That was easy, wasn't it? From now on, we can effortlessly add or remove nodes.

How about upgrading the cluster?

In the next lesson, we will go through the manual upgrading process of the cluster.